

Prediction of the daily mean level of PM10 in Italy in 2020

Github repository [here](#).

General description

The aim of this work is to build machine learning models for predicting the daily mean level of PM10 in Italy in 2020.

Both air quality data and meteorological data will be taken into account.

The data will be processed several times, with different techniques, in order to find the best features and the best dataset. This search is performed with an incremental approach, effectuating a sequence of steps. At each step, several processing alternatives are built and evaluated: the best alternative is then selected and, if his addition on the current dataset does improve the performance, that new feature is added.

Practically speaking, each dataset is evaluated by building a machine learning model on it. For each dataset, four different models are taken into account: each of them is evaluated testing all the possible combinations of some hyperparameters in an exhaustive way. The following list is the list of the four models, with the associated grid of hyperparameters to test.

```
model_paramGrid_list = [  
    ("kNN", KNeighborsRegressor(), {"n_neighbors":range(1,45),  
                                     "weights":["uniform","distance"]}),  
    ("DT", DecisionTreeRegressor(), {"max_leaf_nodes":range(2,51),  
                                     "max_features":["None","sqrt"]}),  
    ("PR", PolynomialRegression(), {"degree":[1,2,3,4]}),  
    ("RF", RandomForestRegressor(), {"n_estimators":range(1,51),  
                                     "max_features":["None","sqrt"]})  
]
```

In the end, the best dataset and the best model are selected (i.e. the best couple dataset-model is selected).

The evaluating measure used during this incremental selection is the MSE cross validation score, which is an error measure: this means that the best couple dataset-model is the one which has associated the minimum validation score.

Prerequisites

The EEA-datasets-handler library will be used in order to acquire and handle the EEA air pollution datasets.

The ILMETEO-datasets-handler library will be used in order to acquire and handle the ILMETEO meteorological datasets.

The timeSeries-processing library will be used in order to process the time series datasets.

Finally, the model-selection library will be used in order to evaluate the machine learning models.

```
>>> import EEA_datasets_handler as eea
>>> import timeSeries_processing as tsp
>>> import model_selection as ms
```

Acquire the EEA air pollution datasets of PM10 in Italy in 2020

```
# Download the datasets
# IT'S NECESSARY ONLY IF THEY HAVEN'T BEEN DOWNLOADED YET
>>> dest_path = "C:\\Datasets"
>>> countries_cities_dict = {"IT": "all"}
>>> pollutants = ["PM10"]
>>> years = [2020]
>>> eea.download_datasets(dest_path, countries_cities_dict,
pollutants, years)

# Load the datasets
>>> source_path = "C:\\Datasets\\EEA"
>>> countries_cities_dict = {"IT": "all"}
>>> pollutants = ["PM10"]
>>> years = [2020]
>>> df = eea.load_datasets(source_path, countries_cities_dict,
pollutants, years)
```

df is a raw DataFrame, which simply contains the measurements: these are not grouped by day and are not cleaned properly.

Acquire the EEA air pollution datasets of PM10 in Italy in all the supported year (i.e. 2013-2020)

```
# Download the datasets
# IT'S NECESSARY ONLY IF THEY HAVEN'T BEEN DOWNLOADED YET
>>> dest_path = "C:\\Datasets"
>>> countries_cities_dict = {"IT": "all"}
>>> pollutants = ["PM10"]
>>> years = [2013, 2014, 2015, 2016, 2017, 2018, 2019, 2020]
>>> eea.download_datasets(dest_path, countries_cities_dict,
pollutants, years)

# Load the datasets
>>> source_path = "C:\\Datasets\\EEA"
>>> countries_cities_dict = {"IT": "all"}
>>> pollutants = ["PM10"]
>>> years = [2013, 2014, 2015, 2016, 2017, 2018, 2019, 2020]
>>> df = eea.load_datasets(source_path, countries_cities_dict,
pollutants, years)
```

df_full is a raw DataFrame, which simply contains the measurements: these are not grouped by day and are not cleaned properly.

Finally, also the meteorological datasets of ILMETEO will be used.

In particular, the meteorological data of the whole Italy are considered, with respect to 2020.

These data are loaded from the local storage ("meteorological_data_2020.csv").

(See the appendix of the chapter about ILMETEO-datasets-handler).

These data will be loaded later on.

Plots

As the plots will show, the evolution over time of the PM10 concentration is strongly periodic.

Plot of the daily mean concentrations of PM10 in Italy in 2020.

```
# Process the 2020 DataFrame (with an arbitrary `fill_n_days`)
>>> df_mean, _, _ = eea.preprocessing(df, fill=True,
fill_n_days=10, fill_aggr="mean")
UserWarning: Missing days: ['2020-01-31', '2020-02-01',
'2020-02-02', '2020-02-03', '2020-02-04', '2020-02-05',
'2020-02-06', '2020-02-07', '2020-02-08', '2020-02-10',
```


```
'2020-02-11']
```

```
>>> df_mean
```

	mean
Datetime	
2020-01-01	76.974569
2020-01-02	56.675791
2020-01-03	55.216906
2020-01-04	54.887035
2020-01-05	28.192059
...	...
2020-12-27	14.997987
2020-12-28	16.317778
2020-12-29	23.536875
2020-12-30	22.759021
2020-12-31	22.005000

```
[366 rows x 1 columns]
```

```
>>> tsp.plot_timeSeries(df_mean, col_name="mean", title="Mean  
concentration of PM10 in Italy in 2020", figsize=(6,6))  
<AxesSubplot:title={'center': 'Mean concentration of PM10 in Italy  
in 2020'}, xlabel='Days', ylabel='mean'>
```




Plot of the daily mean concentrations of PM10 in Italy in all the supported years (2013-2020)

```
# Process the 2013-2020 DataFrame (with an arbitrary
`fill_n_days`)
>>> df_mean_full, _, _ = eea.preprocessing(df_full, fill=True,
fill_n_days=10 ,fill_aggr="mean")
UserWarning: Missing days: ['2020-01-31', '2020-02-01',
'2020-02-02', '2020-02-03', '2020-02-04', '2020-02-05',
'2020-02-06', '2020-02-07', '2020-02-08', '2020-02-10',
'2020-02-11']

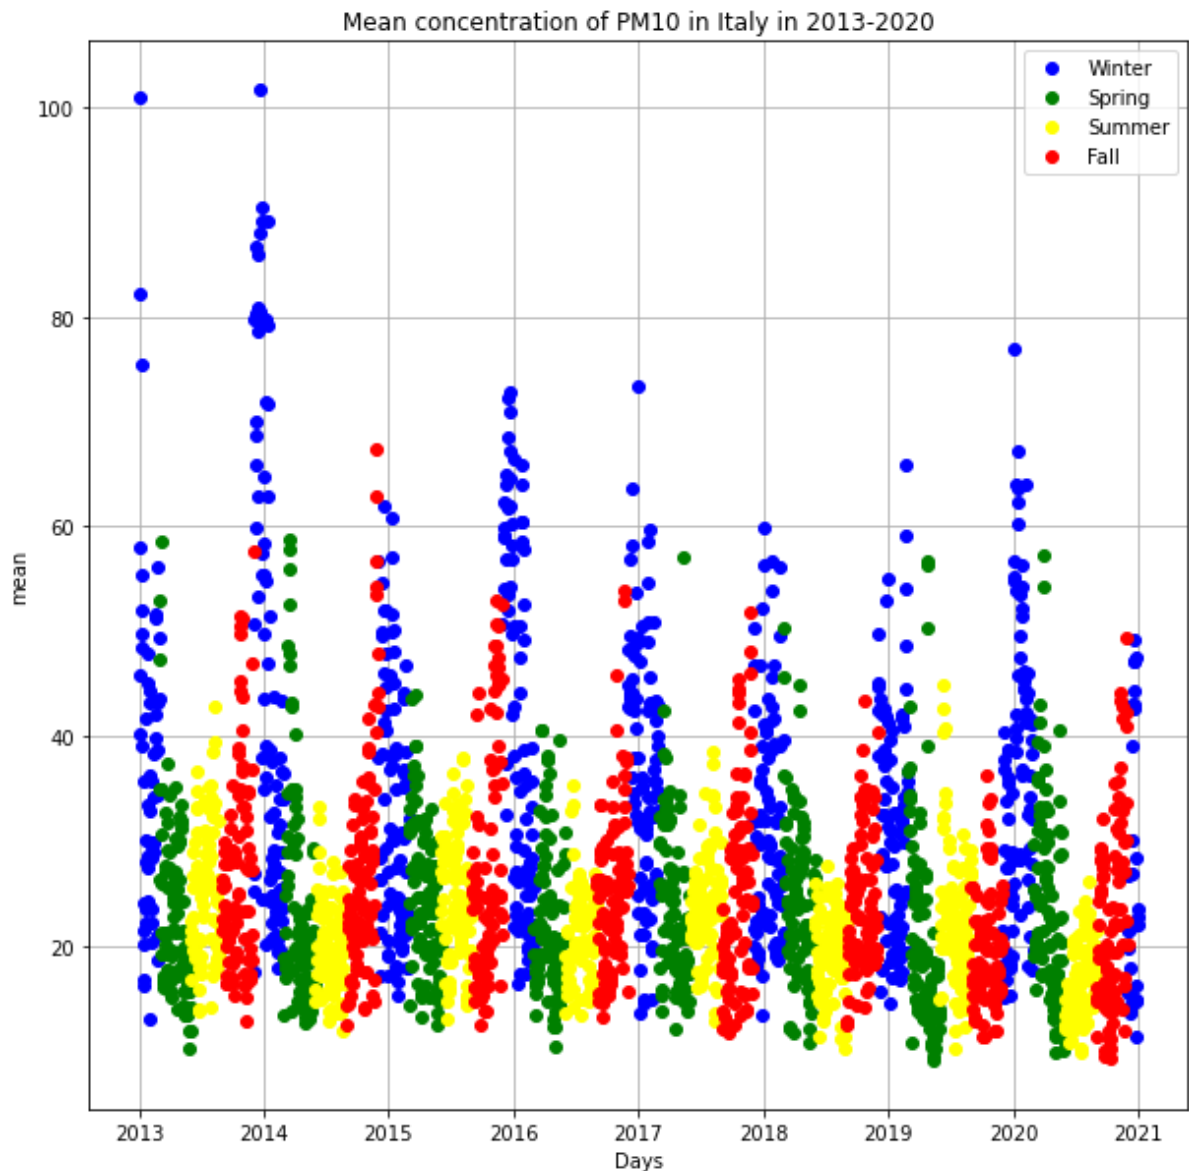
>>> df_mean_full
              mean
Datetime
2013-01-01  100.970568
2013-01-02   40.267997
2013-01-03   45.833017
2013-01-04   58.029136
2013-01-05   82.193141
...          ...
2020-12-27   14.997987
2020-12-28   16.317778
2020-12-29   23.536875
2020-12-30   22.759021
2020-12-31   22.005000

[2922 rows x 1 columns]

>>> tsp.plot_timeSeries(df_mean_full, col_name="mean", title="Mean
concentration of PM10 in Italy in 2013-2020", figsize=(12,12))
<AxesSubplot:title={'center': 'Mean concentration of PM10 in Italy
in 2013-2020'}, xlabel='Days', ylabel='mean'>
```



```
>>> tsp.plot_timeSeries(df_mean_full, col_name="mean",  
divide="season", line=False, title="Mean concentration of PM10 in  
Italy in 2013-2020", figsize=(10,10))  
<AxesSubplot:title={'center': 'Mean concentration of PM10 in Italy  
in 2013-2020'}, xlabel='Days', ylabel='mean'>
```



Plot of the daily mean concentrations of PM10 in Italy during the years 2018-2020

```
>>> df_mean_partial = df_mean_full[df_mean_full.index.map(lambda
day: day.year>=2018)]
```

```
>>> df_mean_partial
      mean
Datetime
2018-01-01  56.360286
2018-01-02  22.650771
2018-01-03  23.279466
```

```

2018-01-04  30.019436
2018-01-05  41.255582
...
2020-12-27  14.997987
2020-12-28  16.317778
2020-12-29  23.536875
2020-12-30  22.759021
2020-12-31  22.005000

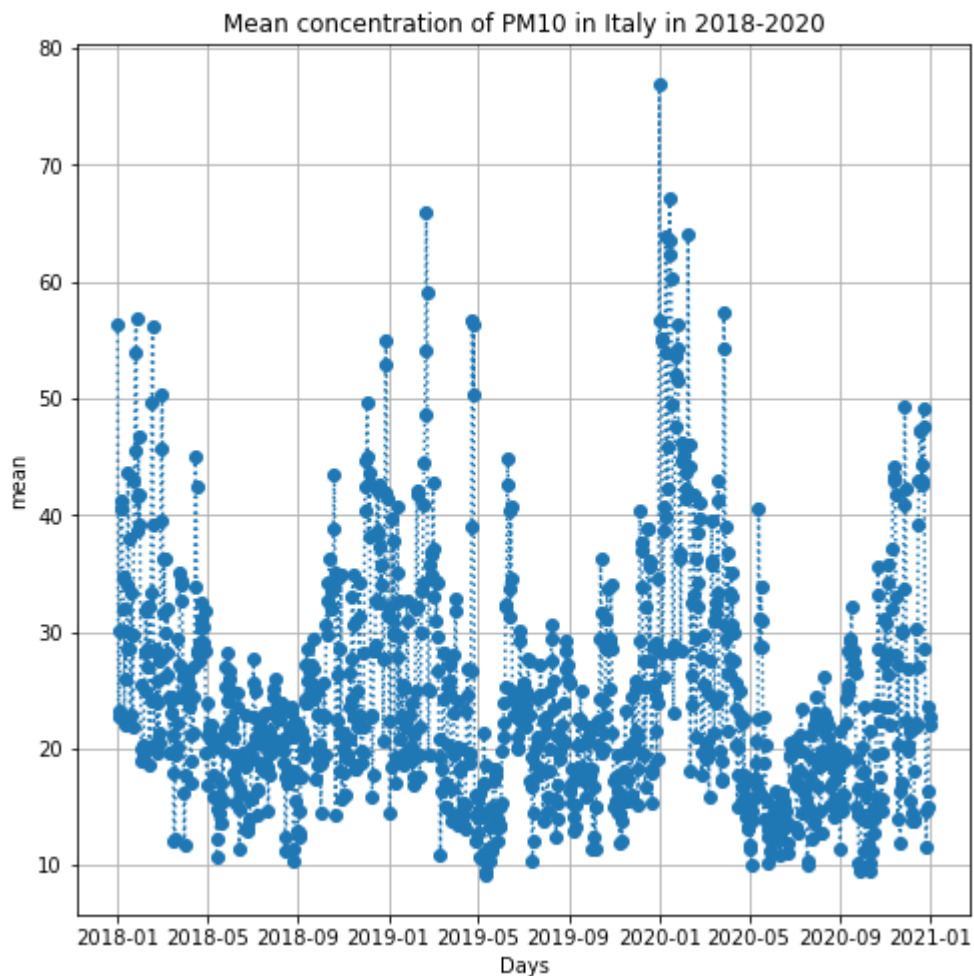
```

```
[2922 rows x 1 columns]
```

```

>>> tsp.plot_timeSeries(df_mean_partial, col_name="mean",
title="Mean concentration of PM10 in Italy in 2018-2020",
figsize=(8,8))
<AxesSubplot:title={'center':'Mean concentration of PM10 in Italy
in 2018-2020'}, xlabel='Days', ylabel='mean'>

```



```

>>> tsp.plot_timeSeries(df_mean_partial, col_name="mean",

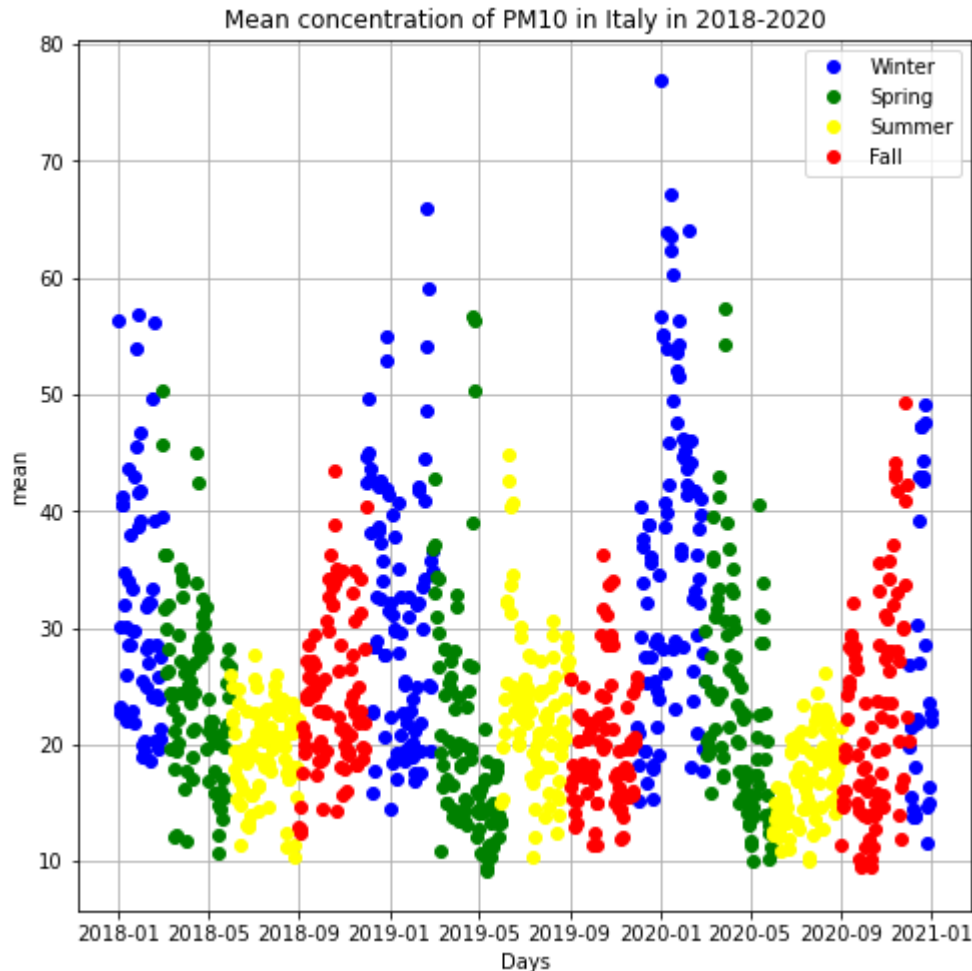
```



```

divide="season", line=False, title="Mean concentration of PM10 in
Italy in 2018-2020", figsize=(8,8))
<AxesSubplot:title={'center':'Mean concentration of PM10 in Italy
in 2018-2020'}, xlabel='Days', ylabel='mean'>

```



Trivial model

Trivial model, which always predicts using the mean of all the daily mean concentrations of PM10 in all the supported years (2013-2020).

```

>>> from sklearn.metrics import mean_squared_error
>>> import numpy as np

# Process the 2013-2020 DataFrame (with an arbitrary
`fill_n_days`)
>>> df_mean_full,_,_ = eea.preprocessing(df_full, fill=True,

```

```

fill_n_days=10, fill_aggr="mean")
UserWarning: Missing days: ['2020-01-31', '2020-02-01',
'2020-02-02', '2020-02-03', '2020-02-04', '2020-02-05',
'2020-02-06', '2020-02-07', '2020-02-08', '2020-02-10',
'2020-02-11']

# Split the DataFrame into X and y: the former contains only the
explanatory features, the latter only the response feature.
# To do that the function `tsp.split_X_y` is used.
>>> X,y = tsp.split_X_y(df_mean, y_col="mean",
                        scale_y=True) # Scale the y

# Splitting the dataset into training-validation-test sets, for
time series, with proportions 0.6-0.2-0.2
>>> train_len = int(X.shape[0]*(0.6))
>>> val_len = int(X.shape[0]*(0.2))
>>> y_val = y[train_len:train_len+val_len] # Validation set

# Trivial predictions
>>> y_pred = np.array([np.mean(y) for i in range(val_len)])

>>> val_score = mean_squared_error(y_true=y_val, y_pred=y_pred)

>>> print("Validation score:", str(val_score))
Validation score: 0.015235327799312809

```

First models

Use k days ago to predict the current day

k new columns, which contain the daily mean level of PM10 in the k preceding days, are added into the 2020 DataFrame.

Two operations have to be made.

1. First of all, the 2020 raw DataFrame df needs to be processed and cleaned into a DataFrame containing the daily mean levels of PM10. The function `eea.processing` will be used. In order to do that, a value for the parameter

`fill_n_days` has to be chosen (i.e. the number of preceding days used to fill a missing day of the 2020 DataFrame).

2. After that, the function `tsp.add_k_previous_days` can be used, in order to create the `k` new features containing the daily mean level of PM10 in the `k` preceding days. To do so, a value for the parameter `k` has to be chosen.

To sum up, two choices have to be made at the same time: the choice for a specific value for `fill_n_days` and the choice for a specific value for `k`.

The value for `fill_n_days` is chosen among the values `[5, 10, 15, "all"]` ("all" means that all the previous days are used to fill a missing day); the value for `k` is chosen among the values `range(3, 15)`. All the possible combinations of the couple `fill_n_days-k` are evaluated, i.e. an exhaustive selection is made. In the end, the best couple is selected.

```
>>> from sklearn.tree import DecisionTreeRegressor
>>> from sklearn.neighbors import KNeighborsRegressor
>>> from sklearn.ensemble import RandomForestRegressor

>>> range_fill_n_days = [5, 10, 15, "all"]
>>> range_k = range(3,15)
# All the datasets that have to be evaluated, one for each
# possible combination of the couple `fill_n_days`-`k`
>>> dataset_list = [
    tsp.add_k_previous_days(eea.preprocessing(df, fill=True,
                                              fill_n_days=fill_n_days,
                                              fill_aggr="mean")[0],
                          col_name="mean", k=k, y_col="mean",
                          scale_y=True)[1:]
    for fill_n_days in range_fill_n_days for k in range_k
]
UserWarning: Missing days: ['2020-01-31', '2020-02-01',
'2020-02-02', '2020-02-03', '2020-02-04', '2020-02-05',
'2020-02-06', '2020-02-07', '2020-02-08', '2020-02-10',
'2020-02-11']
UserWarning: Missing days: ['2020-01-31', '2020-02-01',
'2020-02-02', '2020-02-03', '2020-02-04', '2020-02-05',
'2020-02-06', '2020-02-07', '2020-02-08', '2020-02-10',
'2020-02-11']
UserWarning: Missing days: ['2020-01-31', '2020-02-01',
'2020-02-02', '2020-02-03', '2020-02-04', '2020-02-05',
```

```

'2020-02-06', '2020-02-07', '2020-02-08', '2020-02-10',
'2020-02-11']
...

>>> model_paramGrid_list = [
    ("kNN", KNeighborsRegressor(), {"n_neighbors":range(1,45),
                                    "weights":["uniform","distance"]}),
    ("DT", DecisionTreeRegressor(), {"max_leaf_nodes":range(2,51),
                                    "max_features":[None,"sqrt"]}),
    ("PR", PolynomialRegression(), {"degree":[1,2,3,4]}),
    ("RF", RandomForestRegressor(), {"n_estimators":range(1,51),
                                    "max_features":[None,"sqrt"]})
]
>>> scale_list = [False,True,False,False]

>>> (datasets_train_val_score, datasets_best_model, best_index,
test_score, axes) = ms.datasets_models_validation(dataset_list,
model_paramGrid_list, scale_list=scale_list, time_series=True)

```

Here the results.

```

>>> print("Index best dataset:", str(best_index),
        "; fill_n_days-k:", xvalues[best_index])
Index best dataset: 16 ; fill_n_days-k: 10-7

>>> print("Best model:", datasets_best_model[best_index][1],
        " ; Hyperparameters values:",
        str(datasets_best_model[best_index][2]))
Best model: PR ; Hyperparameters values: {'degree': 1}

>>> print("Validation score:",
        str(datasets_train_val_score[best_index,1]))
Validation score: 0.00646793271790797

>>> print("Test score:", str(test_score))
Test score: 0.01319932892982265

```

Now an example in which the value of fill_n_days has been fixed to 10, in order to show the impact of different values for k on the performance.

```
# In this example only the Polynomial Regression model is used
```

```


>>> range_k = range(3,15)
>>> dataset_list = [
    tsp.add_k_previous_days(eea.preprocessing(df, fill=True,
                                              fill_n_days=10,
                                              fill_aggr="mean")[0],
                           col_name="mean", k=k, y_col="mean",
                           scale_y=True)[1:]

    for k in range_k]
UserWarning: Missing days: ['2020-01-31', '2020-02-01',
'2020-02-02', '2020-02-03', '2020-02-04', '2020-02-05',
'2020-02-06', '2020-02-07', '2020-02-08', '2020-02-10',
'2020-02-11']
UserWarning: Missing days: ['2020-01-31', '2020-02-01',
'2020-02-02', '2020-02-03', '2020-02-04', '2020-02-05',
'2020-02-06', '2020-02-07', '2020-02-08', '2020-02-10',
'2020-02-11']
UserWarning: Missing days: ['2020-01-31', '2020-02-01',
'2020-02-02', '2020-02-03', '2020-02-04', '2020-02-05',
'2020-02-06', '2020-02-07', '2020-02-08', '2020-02-10',
'2020-02-11']
...
>>> xvalues = [str(k) for k in range_k]

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
best_index, test_score,
axes) = ms.datasets_hyperparameter_validation(dataset_list,
    model, hyperparameter, hyperparameter_values, time_series=True,
plot=True, plot_train=True, xvalues=xvalues, xlabel="k",
    title="Datasets validation with different values for k
(fill_n_days=10)")

```



Best dataset: fill_n_days=10 and k=7

There has been an improvement as compared to the trivial model: the validation score has decreased from 0.01523 to 0.00647

```
# Processing with `fill_n_days`=10
>>> df_mean, _, _ = eea.preprocessing(df, fill=True,
fill_n_days=10 ,fill_aggr="mean")
UserWarning: Missing days: ['2020-01-31', '2020-02-01',
'2020-02-02', '2020-02-03', '2020-02-04', '2020-02-05',
'2020-02-06', '2020-02-07', '2020-02-08', '2020-02-10',
'2020-02-11']

# `k`=7
>>> df_mean_best, X, y = tsp.add_k_previous_days(df_mean,
col_name="mean", k=7, y_col="mean", scale_y=True)

>>> df_mean_best
```

	mean	mean_1	mean_2	mean_3	mean_4
Datetime					
2020-01-08	38.646087	40.759729	26.099399	28.192059	54.887035
2020-01-09	39.866008	38.646087	40.759729	26.099399	28.192059
2020-01-10	63.946840	39.866008	38.646087	40.759729	26.099399
2020-01-11	53.902359	63.946840	39.866008	38.646087	40.759729

```

2020-01-12  42.199734  53.902359  63.946840  39.866008  38.646087
...
2020-12-27  14.997987  11.480358  14.695619  28.473515  47.568209
2020-12-28  16.317778  14.997987  11.480358  14.695619  28.473515
2020-12-29  23.536875  16.317778  14.997987  11.480358  14.695619
2020-12-30  22.759021  23.536875  16.317778  14.997987  11.480358
2020-12-31  22.005000  22.759021  23.536875  16.317778  14.997987

      mean_5      mean_6      mean_7
Datetime
2020-01-08  55.216906  56.675791  76.974569
2020-01-09  54.887035  55.216906  56.675791
2020-01-10  28.192059  54.887035  55.216906
2020-01-11  26.099399  28.192059  54.887035
2020-01-12  40.759729  26.099399  28.192059
...
2020-12-27  49.143828  43.004333  42.601218
2020-12-28  47.568209  49.143828  43.004333
2020-12-29  28.473515  47.568209  49.143828
2020-12-30  14.695619  28.473515  47.568209
2020-12-31  11.480358  14.695619  28.473515

[359 rows x 8 columns]

```


Plots which show the goodness of the predictions.

```


>>> model = ms.PolynomialRegression(degree=1) # Best model

>>> ms.plot_predictions(X ,y ,model, xvalues=df_mean_best.index,
xlabel="Days")
<AxesSubplot:title={'center':'Actual vs Predicted values'},
xlabel='Days'>

```



```
>>> ms.plot_predictions(X ,y ,model, plot_type=1,
xvalues=df_mean_best.index, xlabel="Days")
<AxesSubplot:title={'center': 'Actual vs Predicted values'},
xlabel='Actual values', ylabel='Predicted values'>
```



Create the full DataFrame (i.e. years 2013-2020) processed and cleaned (using `fill_n_days=10`).

```
UserWarning: Missing days: ['2020-01-31', '2020-02-01',
'2020-02-02', '2020-02-03', '2020-02-04', '2020-02-05',
'2020-02-06', '2020-02-07', '2020-02-08', '2020-02-10',
'2020-02-11']
```

In this step, two choices have to be made.

- To sum up, two choices have to be made at the same time: the choice for a specific value for `k` and the choice for a specific value for `days_to_select`. All the possible combinations of the couple `k-days_to_select` are evaluated, i.e. an exhaustive selection is made. In the end, the best couple is selected.

[illegible]

[illegible]

```
>>> scale_list = [False, True, False, False]

>>> (datasets_train_val_score, datasets_best_model, best_index,
test_score, axes) = ms.datasets_models_validation(dataset_list,
model_paramGrid_list, scale_list=scale_list, time_series=True)
```

Here the results.

```
>>> print("Index best dataset:", str(best_index),
        "; k-days_to_select:", str(xvalues[best_index]))
Index best dataset: 56 ; k-days_to_select: 6-13

>>> print("Best model:", datasets_best_model[best_index][1],
        "; Hyperparameters values:",
        str(datasets_best_model[best_index][2]))
Best model: PR ; Hyperparameters values: {'degree': 1}

>>> print("Validation score:",
        str(datasets_train_val_score[best_index,1]))
Validation score: 0.005815045651242979

>>> print("Test score:", str(test_score))
Test score: 0.01505297702537988
```

Now an example in which the value of days_to_select has been fixed to 13, in order to show the impact of different values for k on the performance.

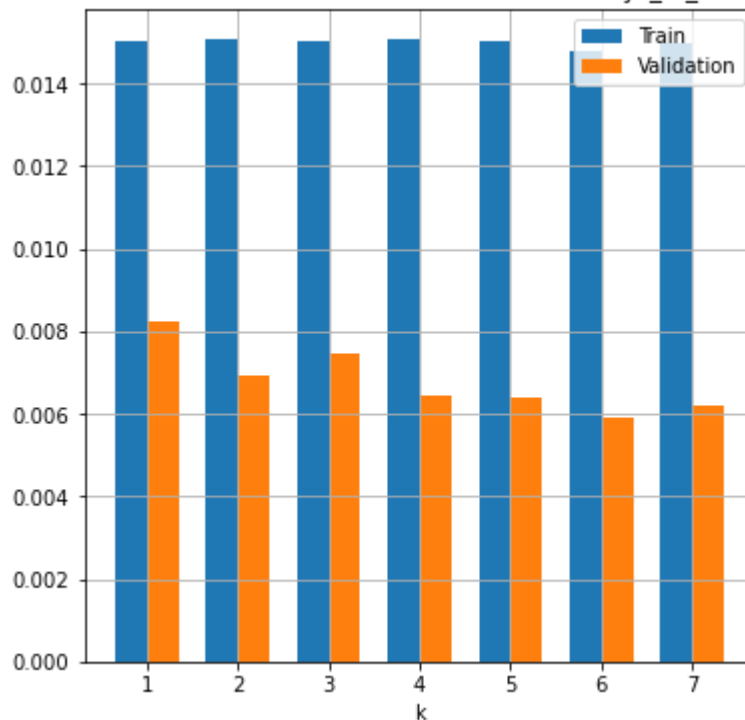
```
# In this example only the Polynomial Regression model is used
>>> range_years = [1, 2, 3, 4, 5, 6, 7]
>>> dataset_list = [
    tsp.add_upTo_k_years_ago_statistics(df=df_mean_best,
                                      df_upTo_k_years_ago=df_full_mean, k=k,
                                      current_year=False,
                                      days_to_select=7, y_col="mean",
                                      scale_y=True)[1:]
    for k in range_years
]
>>> xvalues = [str(k) for k in range_years]

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
```

```
>>> hyperparameter_values=[1, 2, 3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
     best_index, test_score,
     axes) = ms.datasets_hyperparameter_validation(dataset_list,
     model, hyperparameter, hyperparameter_values, time_series=True,
     plot=True, plot_train=True, xvalues=xvalues, xlabel="k",
     title="Datasets validation with differents values for k
     (days_to_select=13)")
```

Datasets validation with differents values for k (days_to_select=13)



Now an example in which the value of k has been fixed to 6, in order to show the impact of different values for days_to_select on the performance.


```
# In this example only the Polynomial Regression model is used
>>> range_n_days = [1, 3, 5, 7, 9, 11, 13, 15, "month", "season"]
>>> dataset_list = [
    tsp.add_upTo_k_years_ago_statistics(df=df_mean_best,
                                      df_upTo_k_years_ago=df_full_mean, k=6,
                                      current_year=False,
                                      days_to_select=days_to_select,
                                      y_col="mean", scale_y=True)[1:]
    for days_to_select in range_n_days
]
```

```
>>> xvalues = [str(days_to_select) for days_to_select in
range_n_days]

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
best_index, test_score,
axes) = ms.datasets_hyperparameter_validation(dataset_list, model,
hyperparameter, hyperparameter_values, time_series=True,
plot=True, plot_train=True, xvalues=xvalues,
xlabel="days_to_select", title="Datasets validation with
differeents values for days_to_select (k=6)")
```

Datasets validation with differents values for days_to_select (k=6)



Best dataset: k=6 and days_to_select=13

There has been an improvement as compared to the previous dataset: the validation score has decreased from 0.00647 to 0.00581.

```
>>> (df_mean_best, _,
_) = tsp.add_upTo_k_years_ago_statistics(df=df_mean_best,
df_upTo_k_years_ago=df_full_mean, k=6, current_year=False,
```

```

days_to_select=13, y_col="mean", scale_y=True)

>>> df_mean_best

```

	mean	mean_1	mean_2	mean_3	mean_4
Datetime					
2020-01-08	38.646087	40.759729	26.099399	28.192059	54.887035
2020-01-09	39.866008	38.646087	40.759729	26.099399	28.192059
2020-01-10	63.946840	39.866008	38.646087	40.759729	26.099399
2020-01-11	53.902359	63.946840	39.866008	38.646087	40.759729
2020-01-12	42.199734	53.902359	63.946840	39.866008	38.646087
...
2020-12-27	14.997987	11.480358	14.695619	28.473515	47.568209
2020-12-28	16.317778	14.997987	11.480358	14.695619	28.473515
2020-12-29	23.536875	16.317778	14.997987	11.480358	14.695619
2020-12-30	22.759021	23.536875	16.317778	14.997987	11.480358
2020-12-31	22.005000	22.759021	23.536875	16.317778	14.997987

	mean_5	mean_6	mean_7	upTo_6_years_ago_mean
Datetime				
2020-01-08	55.216906	56.675791	76.974569	36.287863
2020-01-09	54.887035	55.216906	56.675791	35.814008
2020-01-10	28.192059	54.887035	55.216906	36.143744
2020-01-11	26.099399	28.192059	54.887035	35.806540
2020-01-12	40.759729	26.099399	28.192059	35.522678
...
2020-12-27	49.143828	43.004333	42.601218	38.902204
2020-12-28	47.568209	49.143828	43.004333	38.423588
2020-12-29	28.473515	47.568209	49.143828	37.817480
2020-12-30	14.695619	28.473515	47.568209	36.864552
2020-12-31	11.480358	14.695619	28.473515	35.864452

```

[359 rows x 9 columns]

```

Add statistical information about k years ago

The same full DataFrame prepared before is used, which is `df_mean_full`.

A new column is added, containing statistical information about the daily mean level of PM10 of k years ago. So, the difference with the previous step is that now only a single previous

year is considered. In order to do that, the function `tsp.add_k_years_ago_statistics` is applied.

In this step, two choices have to be made.

1. The specific value of `k` has to be chosen, among the values [1, 2, 3, 4, 5, 6, 7]. (The supported years available are 2013-2020, so up to 7 years ago).
2. The specific value of `days_to_select` has to be chosen, among the values [1, 3, 5, 7, 9, 11, 13, 15, "month", "season"]. This parameter specifies, for each 2020 day, which days of `k` years ago have to be taken into account: if it's a number `n`, the `n` days centered on that day but `k` years ago are selected; if it's "month", the days of the same month but `k` years ago are selected; if it's "season", the days of the same season but `k` years ago are selected.

To sum up, two choices have to be made at the same time: the choice for a specific value for `k` and the choice for a specific value for `days_to_select`. All the possible combinations of the couple `k-days_to_select` are evaluated, i.e. an exhaustive selection is made. In the end, the best couple is selected.

```
>>> from sklearn.tree import DecisionTreeRegressor
>>> from sklearn.neighbors import KNeighborsRegressor
>>> from sklearn.ensemble import RandomForestRegressor

>>> range_years = [1, 2, 3, 4, 5, 6, 7]
>>> range_n_days = [1, 3, 5, 7, 9, 11, 13, 15, "month", "season"]
# All the datasets that have to be evaluated, one for each
# possible combination of the couple `k`-`days_to_select`
>>> dataset_list = [
    tsp.add_k_years_ago_statistics(df_mean_best,
                                  df_k_years_ago=df_full_mean,
                                  k=k,
                                  days_to_select=days_to_select,
                                  y_col="mean",
                                  scale_y=True)[1:]
    for k in range_years for days_to_select in range_n_days
]
```

```
UserWarning: For the day 2020-01-08 only these 7 years ago days have
been found: ['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
'2013-01-05', '2013-01-06', '2013-01-07', '2013-01-08', '2013-01-09',
'2013-01-10', '2013-01-11', '2013-01-12', '2013-01-13', '2013-01-14',
'2013-01-15', '2013-01-16', '2013-01-17', '2013-01-18', '2013-01-19',
'2013-01-20', '2013-01-21', '2013-01-22', '2013-01-23', '2013-01-24',
'2013-01-25', '2013-01-26', '2013-01-27', '2013-01-28', '2013-01-29',
```

```
'2013-01-30', '2013-01-31', '2013-02-01', '2013-02-02', '2013-02-03',
'2013-02-04', '2013-02-05', '2013-02-06', '2013-02-07', '2013-02-08',
'2013-02-09', '2013-02-10', '2013-02-11', '2013-02-12', '2013-02-13',
'2013-02-14', '2013-02-15', '2013-02-16', '2013-02-17', '2013-02-18',
'2013-02-19', '2013-02-20', '2013-02-21', '2013-02-22', '2013-02-23',
'2013-02-24', '2013-02-25', '2013-02-26', '2013-02-27', '2013-02-28']
UserWarning: For the day 2020-01-09 only these 7 years ago days have
been found: ['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
'2013-01-05', '2013-01-06', '2013-01-07', '2013-01-08', '2013-01-09',
'2013-01-10', '2013-01-11', '2013-01-12', '2013-01-13', '2013-01-14',
'2013-01-15', '2013-01-16', '2013-01-17', '2013-01-18', '2013-01-19',
'2013-01-20', '2013-01-21', '2013-01-22', '2013-01-23', '2013-01-24',
'2013-01-25', '2013-01-26', '2013-01-27', '2013-01-28', '2013-01-29',
'2013-01-30', '2013-01-31', '2013-02-01', '2013-02-02', '2013-02-03',
'2013-02-04', '2013-02-05', '2013-02-06', '2013-02-07', '2013-02-08',
'2013-02-09', '2013-02-10', '2013-02-11', '2013-02-12', '2013-02-13',
'2013-02-14', '2013-02-15', '2013-02-16', '2013-02-17', '2013-02-18',
'2013-02-19', '2013-02-20', '2013-02-21', '2013-02-22', '2013-02-23',
'2013-02-24', '2013-02-25', '2013-02-26', '2013-02-27', '2013-02-28']
...
```

```
>>> model_paramGrid_list = [
    ("kNN", KNeighborsRegressor(), {"n_neighbors":range(1,45),
                                     "weights":["uniform","distance"]}),
    ("DT", DecisionTreeRegressor(), {"max_leaf_nodes":range(2,51),
                                     "max_features":[None,"sqrt"]}),
    ("PR", PolynomialRegression(), {"degree":[1,2,3,4]}),
    ("RF", RandomForestRegressor(), {"n_estimators":range(1,51),
                                     "max_features":[None,"sqrt"]})
]
```

```
>>> scale_list = [False,True,False,False]
```

```
>>> (datasets_train_val_score, datasets_best_model, best_index,
test_score, axes) = ms.datasets_models_validation(dataset_list,
model_paramGrid_list, scale_list=scale_list, time_series=True)
```

Here the results.

```
>>> print("Index best dataset:", str(best_index),
          "; k-days_to_select:", str(xvalues[best_index]))
Index best dataset: 50 ; k-days_to_select: 6-1
```



```

>>> print("Best model:", datasets_best_model[best_index][1],
        "; Hyperparameters values:",
        str(datasets_best_model[best_index][2]))
Best model: PR ; Hyperparameters values: {'degree': 1}

>>> print("Validation score:",
        str(datasets_train_val_score[best_index,1]))
Validation score: 0.005804228658914929

>>> print("Test score:", str(test_score))
Test score: 0.014628394398884796

```

Now an example in which the value of `days_to_select` has been fixed to 1, in order to show the impact of different values for `k` on the performance.

```

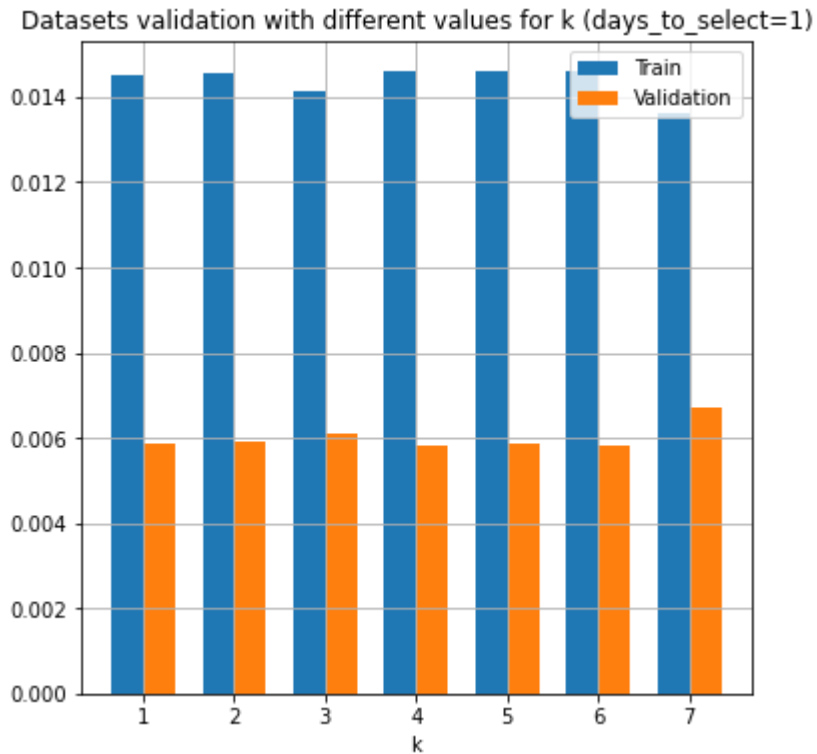
# In this example only the Polynomial Regression model is used
>>> range_years = [1, 2, 3, 4, 5, 6, 7]
>>> dataset_list = [
    tsp.add_k_years_ago_statistics(df_mean_best,
                                  df_k_years_ago=df_full_mean,
                                  k=k, days_to_select=1,
                                  y_col="mean",
                                  scale_y=True)[1:]

    for k in range_years]
>>> xvalues = [str(k) for k in range_years]

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1, 2, 3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
    best_index, test_score,
    axes) = ms.datasets_hyperparameter_validation(dataset_list, model,
    hyperparameter, hyperparameter_values, time_series=True,
    plot=True, plot_train=True, xvalues=xvalues, xlabel="k",
    title="Datasets validation with different values for k
    (days_to_select=1)")

```



Now an example in which the value of k has been fixed to 6, in order to show the impact of different values for days_to_select on the performance.


```
# In this example only the Polynomial Regression model is used
>>> range_n_days = [1, 3, 5, 7, 9, 11, 13, 15, "month", "season"]
>>> dataset_list = [
    tsp.add_k_years_ago_statistics(df_mean_best,
                                  df_k_years_ago=df_full_mean,
                                  k=6,
                                  days_to_select=days_to_select,
                                  y_col="mean",
                                  scale_y=True)[1:]

    for days_to_select in range_n_days]
>>> xvalues = [str(days_to_select) for days_to_select in
range_n_days]

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
    best_index, test_score,
    axes) = ms.datasets_hyperparameter_validation(dataset_list, model,
```

```
hyperparameter, hyperparameter_values, time_series=True,
plot=True, plot_train=True, xvalues=xvalues,
xlabel="days_to_select", title="Datasets validation with different
values for days_to_select (k=6)")
```



Best dataset: k=6 and days_to_select=1

There has been an improvement as compared to the previous dataset: the validation score has decreased from 0.00581 to 0.00580.

```
>>> (df_mean_best, _,
_) = tsp.add_k_years_ago_statistics(df_mean_best,
df_k_years_ago=df_full_mean, k=6, days_to_select=1, y_col="mean",
scale_y=True)
```

```
>>> df_mean_best
```

	mean	mean_1	mean_2	mean_3	mean_4
Datetime					
2020-01-08	38.646087	40.759729	26.099399	28.192059	54.887035
2020-01-09	39.866008	38.646087	40.759729	26.099399	28.192059
2020-01-10	63.946840	39.866008	38.646087	40.759729	26.099399
2020-01-11	53.902359	63.946840	39.866008	38.646087	40.759729
2020-01-12	42.199734	53.902359	63.946840	39.866008	38.646087

```

...
2020-12-27 14.997987 11.480358 14.695619 28.473515 47.568209
2020-12-28 16.317778 14.997987 11.480358 14.695619 28.473515
2020-12-29 23.536875 16.317778 14.997987 11.480358 14.695619
2020-12-30 22.759021 23.536875 16.317778 14.997987 11.480358
2020-12-31 22.005000 22.759021 23.536875 16.317778 14.997987

```

```

          mean_5      mean_6      mean_7  upTo_6_years_ago_mean
Datetime
2020-01-08 55.216906  56.675791  76.974569      36.287863
2020-01-09 54.887035  55.216906  56.675791      35.814008
2020-01-10 28.192059  54.887035  55.216906      36.143744
2020-01-11 26.099399  28.192059  54.887035      35.806540
2020-01-12 40.759729  26.099399  28.192059      35.522678
...
2020-12-27 49.143828  43.004333  42.601218      38.902204
2020-12-28 47.568209  49.143828  43.004333      38.423588
2020-12-29 28.473515  47.568209  49.143828      37.817480
2020-12-30 14.695619  28.473515  47.568209      36.864552
2020-12-31 11.480358  14.695619  28.473515      35.864452

```

```

          6_years_ago_mean
Datetime
2020-01-08      71.802620
2020-01-09      79.803419
2020-01-10      71.632078
2020-01-11      79.216124
2020-01-12      89.111978
...
2020-12-27      17.186082
2020-12-28      19.039899
2020-12-29      20.643419
2020-12-30      19.247014
2020-12-31      23.937203

```

```
[359 rows x 10 columns]
```

Add meteorological information

Load the meteorological dataset, which contains values for several climatic agents in Italy in 2020.

```
>>> import pandas as pd
>>> df_meteo = pd.read_csv("meteorological_data_2020.csv",
index_col=0)
>>> df_meteo = df_meteo.set_index(pd.DatetimeIndex(df_meteo.index))

>>> df_meteo
```

	TMEDIA °C	TMIN °C	TMAX °C	PUNTORUGIADA °C
2020-01-01	5.202196	0.173764	11.003131	2.534921
2020-01-02	4.774801	0.182640	10.325979	2.465371
2020-01-03	4.885063	0.531284	9.658314	3.231408
2020-01-04	6.162960	1.381113	11.354987	3.776887
2020-01-05	6.551925	1.686219	12.409042	3.172703
...
2020-12-27	3.366679	0.191966	6.335378	2.162273
2020-12-28	4.014919	1.075045	6.692360	3.125000
2020-12-29	4.921844	1.494424	8.174234	3.427775
2020-12-30	5.007944	2.084797	8.026350	3.539751
2020-12-31	3.614142	-0.334781	7.465027	2.430201
...				

Add meteorological information about 2020

A new column is added, containing meteorological information about a specific climatic agent in Italy in 2020.

That specific climatic agent has to be chosen among the list ["TMEDIA °C", "TMIN °C", "TMAX °C", "UMIDITA %", "VENTOMEDIA km/h", "VENTOMAX km/h", "PRESSIONESLM mb"]. All these possible agents are evaluated: in the end, the best one is selected.

In order to add a climatic agent, the function `tsp.add_timeSeries_dataframe` is used.

```
>>> from sklearn.tree import DecisionTreeRegressor
>>> from sklearn.neighbors import KNeighborsRegressor
```

```

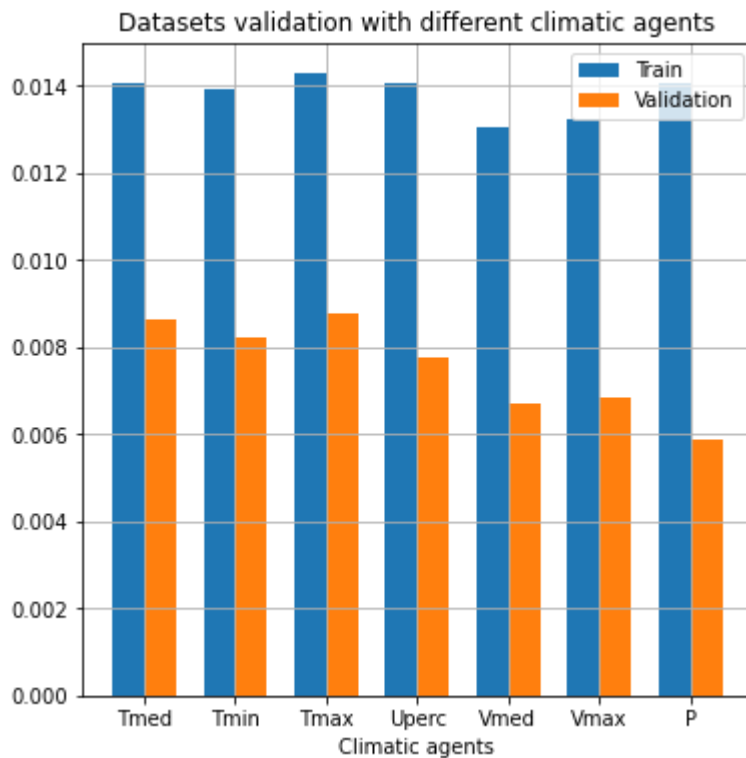
>>> from sklearn.ensemble import RandomForestRegressor

>>> columns = ["TMEDIA °C", "TMIN °C", "TMAX °C", "UMIDITA %",
"VENTOMEDIA km/h", "VENTOMAX km/h", "PRESSIONESLM mb"]
# All the datasets that have to be evaluated, one for each
possible climatic agent
>>> dataset_list = [
    tsp.add_timeSeries_dataframe(df_mean_best, df_meteo[[agent]],
                                y_col="mean", scale_y=True)[1:]
    for agent in columns]
>>> xvalues = ["Tmed", "Tmin", "Tmax", "Uperc", "Vmed", "Vmax",
"P"]

>>> model_paramGrid_list = [
    ("kNN", KNeighborsRegressor(), {"n_neighbors":range(1,45),
                                    "weights":["uniform","distance"]}),
    ("DT", DecisionTreeRegressor(), {"max_leaf_nodes":range(2,51),
                                     "max_features":[None,"sqrt"]}),
    ("PR", PolynomialRegression(), {"degree":[1,2,3,4]}),
    ("RF", RandomForestRegressor(), {"n_estimators":range(1,51),
                                     "max_features":[None,"sqrt"]})
]
>>> scale_list = [False, True, False, False]

>>> (datasets_train_val_score, datasets_best_model, best_index,
test_score, axes) = ms.datasets_models_validation(dataset_list,
model_paramGrid_list, scale_list=scale_list, time_series=True,
regr=True, plot=True, plot_train=True, xvalues=xvalues,
xlabel="Climatic agents", title="Datasets validation with
different climatic agents")

```



Here the results.

```
>>> print("Index best dataset:", str(best_index),
          "; Climatic agent:", str(xvalues[best_index]))
Index best dataset: 6 ; Climatic agent: P

>>> print("Best model:", datasets_best_model[best_index][1],
          "; Hyperparameters values:",
          str(datasets_best_model[best_index][2]))
Best model: PR ; Hyperparameters values: {'degree': 1}

>>> print("Validation score:",
          str(datasets_train_val_score[best_index,1]))
Validation score: 0.005893619946527116

>>> print("Test score:", str(test_score))
Test score: 0.011106647607909125
```

There has been no improvement as compared to the previous dataset: no new feature is added.

The validation score associated with the best climatic agent has not decreased as compared to the previous dataset: the previous validation score was 0.00580 and the current is 0.00589

In all the datasets that have been tested, the PolynomialRegression model was always the best one. For this reason, from this point on only this model will be used.

Add statistical information about 2020 days with similar meteorological conditions

The previous attempt to add meteorological features was a failure. For this reason, the technique and the point of view have to be changed.

First of all, a new DataFrame is created: this simply contains the daily mean levels of PM10 in Italy in 2020 and the climatic agents in Italy in 2020. This DataFrame is simply obtained by merging the DataFrames df_mean and df_meteo.

```
>>> df_mean_meteo, _, _ = tsp.add_timeSeries_dataframe(df_mean,
df_meteo, y_col="mean", scale_y=True)
>>> df_mean_meteo
```

	mean	TMEDIA °C	TMIN °C	TMAX °C
2020-01-01	76.974569	5.202196	0.173764	11.003131
2020-01-02	56.675791	4.774801	0.182640	10.325979
2020-01-03	55.216906	4.885063	0.531284	9.658314
2020-01-04	54.887035	6.162960	1.381113	11.354987
2020-01-05	28.192059	6.551925	1.686219	12.409042
...
2020-12-27	14.997987	3.366679	0.191966	6.335378
2020-12-28	16.317778	4.014919	1.075045	6.692360
2020-12-29	23.536875	4.921844	1.494424	8.174234
2020-12-30	22.759021	5.007944	2.084797	8.026350
2020-12-31	22.005000	3.614142	-0.334781	7.465027
...

A new column is added, containing statistical information about the daily mean level of PM10 of the preceding 2020 days with similar meteorological conditions. So, for each 2020 day the

preceding 2020 days with similar meteorological conditions are considered and then the mean of their daily level of PM10 is computed.

In order to do that, the function `tsp.add_current_year_statistics` is applied, specifying as `days_to_select` a predicate (i.e. a function which returns a boolean).

Now, the meaning of 'similar meteorological conditions' has to be defined. Two days have similar meteorological conditions if, for a certain specified climatic agent, their values differ by not more than a certain tolerance `t`. So, two choices have to be made:

- the climatic agent has to be chosen, among ["TMEDIA °C", "TMIN °C", "TMAX °C", "UMIDITA %", "VENTOMEDIA km/h", "VENTOMAX km/h", "PRESSIONESLM mb"]`
- the tolerance has to be chosen (the possible values depend on the climatic agent).

All the possible combinations of the couple `Climatic agent-t` are evaluated, i.e. an exhaustive selection is made. In the end, the best couple is selected. Unlike the dataset selections seen previously, now the dataset selection is divided by climatic agent, in order to simplify the visualization.

As it will be shown, the just described technique is effective. In fact, four new features will be added with this technique.

First addition

TMEDIA

```
# Higher-order function, which creates the `days_to_select`
function.
# It takes in input a specific tolerance `t` and it returns the
corresponding `days_to_select` function.
>>> def create_daysToSelect_function(t):
    # `days_to_select` is a lambda function which decides whether
    # to select `day_current_year` for `day` or not.
    # It returns True if `day_current_year` has a similar TMEDIA
    # with respect to `day` (i.e. the difference is less than `t`)
    return (
        lambda day, df, day_current_year, df_current_year :
            abs(df_current_year["TMEDIA °C"].loc[day] -
              df_current_year["TMEDIA °C"].loc[day_current_year])<t
    )

>>> tolerances = [0.5, 1, 2, 4, 6, 8, 10] # Possible tolerances
>>> dataset_list = [
```

```


tsp.add_current_year_statistics(df=df_mean_best,
                               df_current_year=df_mean_meteo,
                               # Create `days_to_select`
                               days_to_select = create_daysToSelect_function(t),
                               # Only the "mean" column is of interest
                               columns_to_select=["mean"], y_col="mean",
                               scale_y=True)[1:]
for t in tolerances
]
UserWarning: No current year days have been found for the day
2020-01-08
UserWarning: No current year days have been found for the day
2020-01-28
UserWarning: No current year days have been found for the day
2020-02-03
UserWarning: No current year days have been found for the day
2020-02-04
...
>>> xvalues = [str(t) for t in tolerances]

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
    best_index, test_score,
    axes) = ms.datasets_hyperparameter_validation(dataset_list, model,
    hyperparameter, hyperparameter_values, time_series=True,
    plot=True, plot_train=True, xvalues=xvalues, xlabel="Tolerances",
    title="Datasets validation with respect to TMEDIA using different
    tolerances")

```

Datasets validation with respect to TMEDIA using different tolerances



```
>>> print("Index best dataset:", str(best_index), "; Tolerance:",
          str(xvalues[best_index]))
Index best dataset: 1 ; Tolerance: 1

>>> print("Best polynomial grade:",
          str(datasets_best_hyperparameter_value[best_index]))
Best polynomial grade: 1

>>> print("Validation score:",
          str(datasets_train_val_score[best_index,1]))
Validation score: 0.005973842073088957

>>> print("Test score:", str(test_score))
Test score: 0.014607561953069621
```

The validation score has not decreased.

TMIN

```
# Higher-order function, which creates the `days_to_select`
function.
# It takes in input a specific tolerance `t` and it returns the
corresponding `days_to_select` function.
```

```

>>> def create_daysToSelect_function(t):
    # `days_to_select` is a lambda function which decides whether
    # to select `day_current_year` for `day` or not.
    # It returns True if `day_current_year` has a similar TMIN
    # with respect to `day` (i.e. the difference is less than `t`)
    return (
        lambda day, df, day_current_year, df_current_year :
            abs(df_current_year["TMIN °C"].loc[day] -
                df_current_year["TMIN °C"].loc[day_current_year])<t
    )

>>> tolerances = [0.5, 1, 2, 4, 6, 8, 10] # Possible tolerances
>>> dataset_list = [
    tsp.add_current_year_statistics(df=df_mean_best,
                                   df_current_year=df_mean_meteo,
                                   # Create `days_to_select`
                                   days_to_select = create_daysToSelect_function(t),
                                   # Only the "mean" column is of interest
                                   columns_to_select=["mean"], y_col="mean",
                                   scale_y=True)[1:]
    for t in tolerances
]
UserWarning: No current year days have been found for the day
2020-01-08
UserWarning: No current year days have been found for the day
2020-01-18
UserWarning: No current year days have been found for the day
2020-01-25
UserWarning: No current year days have been found for the day
2020-01-28
...
>>> xvalues = [str(t) for t in tolerances]


>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
    best_index, test_score,
    axes) = ms.datasets_hyperparameter_validation(dataset_list,
    model, hyperparameter, hyperparameter_values, time_series=True,

```

```
plot=True, plot_train=True, xvalues=xvalues, xlabel="Tolerances",
title="Datasets validation with respect to TMIN using different
tolerances")
```

Datasets validation with respect to TMIN using different tolerances



```
>>> print("Index best dataset:", str(best_index), "; Tolerance:",
          str(xvalues[best_index]))
Index best dataset: 2 ; Tolerance: 2

>>> print("Best polynomial grade:",
          str(datasets_best_hyperparameter_value[best_index]))
Best polynomial grade: 1

>>> print("Validation score:",
          str(datasets_train_val_score[best_index,1]))
Validation score: 0.00585275308627106

>>> print("Test score:", str(test_score))
Test score: 0.013347518888114866
```

The validation score has not decreased.

TMAX

```

# Higher-order function, which creates the `days_to_select`
function.
# It takes in input a specific tolerance `t` and it returns the
corresponding `days_to_select` function.
>>> def create_daysToSelect_function(t):
    # `days_to_select` is a lambda function which decides whether
    # to select `day_current_year` for `day` or not.
    # It returns True if `day_current_year` has a similar TMAX
    # with respect to `day` (i.e. the difference is less than `t`)
    return (
        lambda day, df, day_current_year, df_current_year :
            abs(df_current_year["TMAX °C"].loc[day] -
                df_current_year["TMAX °C"].loc[day_current_year])<t
    )


>>> tolerances = [0.5, 1, 2, 4, 6, 8, 10] # Possible tolerances
>>> dataset_list = [
    tsp.add_current_year_statistics(df=df_mean_best,
                                   df_current_year=df_mean_meteo,
                                   # Create `days_to_select`
                                   days_to_select = create_daysToSelect_function(t),
                                   # Only the "mean" column is of interest
                                   columns_to_select=["mean"], y_col="mean",
                                   scale_y=True)[1:]
    for t in tolerances
]
UserWarning: No current year days have been found for the day
2020-01-29
UserWarning: No current year days have been found for the day
2020-01-31
UserWarning: No current year days have been found for the day
2020-02-03
UserWarning: No current year days have been found for the day
2020-02-04
...
>>> xvalues = [str(t) for t in tolerances]

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

```

```
>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
      best_index, test_score,
      axes) = ms.datasets_hyperparameter_validation(dataset_list,
      model, hyperparameter, hyperparameter_values, time_series=True,
      plot=True, plot_train=True, xvalues=xvalues, xlabel="Tolerances",
      title="Datasets validation with respect to TMAX using different
      tolerances")
```

Datasets validation with respect to TMAX using different tolerances



```
>>> print("Index best dataset:", str(best_index), "; Tolerance:",
          str(xvalues[best_index]))
Index best dataset: 2 ; Tolerance: 2

>>> print("Best polynomial grade:",
          str(datasets_best_hyperparameter_value[best_index]))
Best polynomial grade: 1

>>> print("Validation score:",
          str(datasets_train_val_score[best_index,1]))
Validation score: 0.0062306936890401425

>>> print("Test score:", str(test_score))
Test score: 0.015206419159237023
```

The validation score has not decreased.

VENTOMEDIA

```
# Higher-order function, which creates the `days_to_select`
function.
# It takes in input a specific tolerance `t` and it returns the
corresponding `days_to_select` function.
>>> def create_daysToSelect_function(t):
    # `days_to_select` is a lambda function which decides whether
    # to select `day_current_year` for `day` or not.
    # It returns True if `day_current_year` has a similar
    # VENTOMEDIA with respect to `day` (i.e. the difference is
    # less than `t`)
    return (
        lambda day, df, day_current_year, df_current_year :
            abs(df_current_year["VENTOMEDIA km/h"].loc[day] -
                df_current_year["VENTOMEDIA km/h"].loc[day_current_year])<t
    )


>>> tolerances = [0.3, 0.5, 1, 2, 3, 5, 10, 15] # Possible
tolerances
>>> dataset_list = [
    tsp.add_current_year_statistics(df=df_mean_best,
                                   df_current_year=df_mean_meteo,
                                   # Create `days_to_select`
                                   days_to_select = create_daysToSelect_function(t),
                                   # Only the "mean" column is of interest
                                   columns_to_select=["mean"], y_col="mean",
                                   scale_y=True)[1:]
    for t in tolerances
]
UserWarning: No current year days have been found for the day
2020-01-09
UserWarning: No current year days have been found for the day
2020-01-18
UserWarning: No current year days have been found for the day
2020-01-19
UserWarning: No current year days have been found for the day
2020-01-20
...
>>> xvalues = [str(t) for t in tolerances]
```



```
>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
    best_index, test_score,
    axes) = ms.datasets_hyperparameter_validation(dataset_list,
    model, hyperparameter, hyperparameter_values, time_series=True,
    plot=True, plot_train=True, xvalues=xvalues, xlabel="Tolerances",
    title="Datasets validation with respect to VENTOMEDIA using
    different tolerances")
```

Datasets validation with respect to VENTOMEDIA using different tolerances



```
>>> print("Index best dataset:", str(best_index), "; Tolerance:",
        str(xvalues[best_index]))
Index best dataset: 2 ; Tolerance: 1

>>> print("Best polynomial grade:",
        str(datasets_best_hyperparameter_value[best_index]))
Best polynomial grade: 1

>>> print("Validation score:",
```

```

        str(datasets_train_val_score[best_index,1]))
Validation score: 0.00557915970824468

>>> print("Test score:", str(test_score))
Test score: 0.0070760470805857805

```

The validation score has decreased!

VENTOMAX

```

# Higher-order function, which creates the `days_to_select`
function.
# It takes in input a specific tolerance `t` and it returns the
corresponding `days_to_select` function.
>>> def create_daysToSelect_function(t):
    # `days_to_select` is a lambda function which decides whether
    # to select `day_current_year` for `day` or not.
    # It returns True if `day_current_year` has a similar
    # VENTOMAX with respect to `day` (i.e. the difference is
    # less than `t`)
    return (
        lambda day, df, day_current_year, df_current_year :
            abs(df_current_year["VENTOMAX km/h"].loc[day] -
                df_current_year["VENTOMAX km/h"].loc[day_current_year])<t
    )

>>> tolerances = [0.2, 0.5, 1, 2, 3, 5, 10, 15] # Possible
tolerances
>>> dataset_list = [
    tsp.add_current_year_statistics(df=df_mean_best,
                                   df_current_year=df_mean_meteo,
                                   # Create `days_to_select`
                                   days_to_select = create_daysToSelect_function(t),
                                   # Only the "mean" column is of interest
                                   columns_to_select=["mean"], y_col="mean",
                                   scale_y=True)[1:]
    for t in tolerances
]
UserWarning: No current year days have been found for the day
2020-01-09
UserWarning: No current year days have been found for the day
2020-01-13
UserWarning: No current year days have been found for the day

```

2020-01-18

UserWarning: No current year days have been found for the day

2020-01-19

...

```
>>> xvalues = [str(t) for t in tolerances]
```


```
>>> model = ms.PolynomialRegression()
```

```
>>> hyperparameter="degree"
```

```
>>> hyperparameter_values=[1,2,3]
```

```
>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
    best_index, test_score,
    axes) = ms.datasets_hyperparameter_validation(dataset_list,
    model, hyperparameter, hyperparameter_values, time_series=True,
    plot=True, plot_train=True, xvalues=xvalues, xlabel="Tolerances",
    title="Datasets validation with respect to VENTOMAX using
    different tolerances")
```

Datasets validation with respect to VENTOMAX using different tolerances



```
>>> print("Index best dataset:", str(best_index), "; Tolerance:",
    str(xvalues[best_index]))
```

Index best dataset: 4 ; Tolerance: 3

```
>>> print("Best polynomial grade:",
```

```

        str(datasets_best_hyperparameter_value[best_index]))
Best polynomial grade: 1

>>> print("Validation score:",
        str(datasets_train_val_score[best_index,1]))
Validation score: 0.005414068416456359

>>> print("Test score:", str(test_score))
Test score: 0.0071246954594039215

```

The validation score has decreased!

PRESSIONE

```

# Higher-order function, which creates the `days_to_select`
function.
# It takes in input a specific tolerance `t` and it returns the
corresponding `days_to_select` function.
>>> def create_daysToSelect_function(t):
    # `days_to_select` is a lambda function which decides whether
    # to select `day_current_year` for `day` or not.
    # It returns True if `day_current_year` has a similar
    # PRESSIONE with respect to `day` (i.e. the difference is
    # less than `t`)
    return (
        lambda day, df, day_current_year, df_current_year :
            abs(df_current_year["PRESSIONESLM mb"].loc[day] -
                df_current_year["PRESSIONESLM mb"].loc[day_current_year])<t
    )

>>> tolerances = [0.2, 0.5, 1, 3, 5, 10, 15, 20] # Possible
tolerances
>>> dataset_list = [
    tsp.add_current_year_statistics(df=df_mean_best,
                                   df_current_year=df_mean_meteo,
                                   # Create `days_to_select`
                                   days_to_select = create_daysToSelect_function(t),
                                   # Only the "mean" column is of interest
                                   columns_to_select=["mean"], y_col="mean",
                                   scale_y=True)[1:]
    for t in tolerances
]

```

```


UserWarning: No current year days have been found for the day
2020-01-09
UserWarning: No current year days have been found for the day
2020-01-10
UserWarning: No current year days have been found for the day
2020-01-11
UserWarning: No current year days have been found for the day
2020-01-13
...
>>> xvalues = [str(t) for t in tolerances]

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
    best_index, test_score,
    axes) = ms.datasets_hyperparameter_validation(dataset_list,
    model, hyperparameter, hyperparameter_values, time_series=True,
    plot=True, plot_train=True, xvalues=xvalues, xlabel="Tolerances",
    title="Datasets validation with respect to PRESSIONESLM using
    different tolerances")

```

Datasets validation with respect to PRESSIONESLM using different tolerances



```

>>> print("Index best dataset:", str(best_index), "; Tolerance:",
          str(xvalues[best_index]))
Index best dataset: 0 ; Tolerance: 0.2

>>> print("Best polynomial grade:",
          str(datasets_best_hyperparameter_value[best_index]))
Best polynomial grade: 1

>>> print("Validation score:",
          str(datasets_train_val_score[best_index,1]))
Validation score: 0.005665634068097989

>>> print("Test score:", str(test_score))
Test score: 0.01260854713712688

```

The validation score has decreased!

UMIDITA

```

# Higher-order function, which creates the `days_to_select`
function.
# It takes in input a specific tolerance `t` and it returns the
corresponding `days_to_select` function.
>>> def create_daysToSelect_function(t):
    # `days_to_select` is a lambda function which decides whether
    # to select `day_current_year` for `day` or not.
    # It returns True if `day_current_year` has a similar
    # UMIDITA with respect to `day` (i.e. the difference is
    # less than `t`)
    return (
        lambda day, df, day_current_year, df_current_year :
            abs(df_current_year["UMIDITA %"].loc[day] -
               df_current_year["UMIDITA %"].loc[day_current_year])<t
    )

>>> tolerances = [0.4, 0.7, 1, 2, 4, 6, 8, 10] # Possible
tolerances
>>> dataset_list = [
    tsp.add_current_year_statistics(df=df_mean_best,
                                   df_current_year=df_mean_meteo,
                                   # Create `days_to_select`
                                   days_to_select = create_daysToSelect_function(t),

```

```


        # Only the "mean" column is of interest
        columns_to_select=["mean"], y_col="mean",
        scale_y=True)[1:]
    for t in tolerances
]
UserWarning: No current year days have been found for the day
2020-01-10
UserWarning: No current year days have been found for the day
2020-01-11
UserWarning: No current year days have been found for the day
2020-01-12
UserWarning: No current year days have been found for the day
2020-01-14
...
>>> xvalues = [str(t) for t in tolerances]

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
    best_index, test_score,
    axes) = ms.datasets_hyperparameter_validation(dataset_list,
    model, hyperparameter, hyperparameter_values, time_series=True,
    plot=True, plot_train=True, xvalues=xvalues, xlabel="Tolerances",
    title="Datasets validation with respect to UMIDITA using different
    tolerances")

```

Datasets validation with respect to UMIDITA using different tolerances



```
>>> print("Index best dataset:", str(best_index), "; Tolerance:",
          str(xvalues[best_index]))
Index best dataset: 2 ; Tolerance: 1

>>> print("Best polynomial grade:",
          str(datasets_best_hyperparameter_value[best_index]))
Best polynomial grade: 1

>>> print("Validation score:",
          str(datasets_train_val_score[best_index,1]))
Validation score: 0.005827590144634284

>>> print("Test score:", str(test_score))
Test score: 0.014934800194941403
```

The validation score has not decreased

VMEDIA VMAX PRESSIONE

Trying to use more climatic agents together could be interesting. The three most effective climatic agents are considered, which are VMEDIA, VMAX, P. Each of them is tested with different tolerances: t_1 is the tolerance for VMEDIA, t_2 is the tolerance for VMAX, and t_3 is the tolerance for P. All the possible combinations of tolerances $t_1-t_2-t_3$ are evaluated.


```

# Higher-order function, which creates the `days_to_select` function.
# It takes in input three specific tolerances `t1-t2-t3` and it returns
the corresponding `days_to_select` function.
>>> def create_daysToSelect_function(t1, t2, t3):
    return (
        lambda day, df, day_current_year, df_current_year :
        (abs(df_current_year["VENTOMEDIA km/h"].loc[day] -
            df_current_year["VENTOMEDIA km/h"].loc[day_current_year])<t1 and
            abs(df_current_year["VENTOMAX km/h"].loc[day] -
            df_current_year["VENTOMAX km/h"].loc[day_current_year])<t2 and
            abs(df_current_year["PRESSIONESLM mb"].loc[day] -
            df_current_year["PRESSIONESLM mb"].loc[day_current_year])<t3)
    )

>>> tolerances = [1, 5, 10, 15] # Possible tolerances
>>> dataset_list = [
    tsp.add_current_year_statistics(df=df_mean_best,
        df_current_year=df_mean_meteo,
        # Create `days_to_select`
        days_to_select=create_daysToSelect_function(t1,t2,t3),
        # Only the "mean" column is of interest
        columns_to_select=["mean"], y_col="mean",
        scale_y=True)[1:]
    for t1 in tolerances for t2 in tolerances for t3 in tolerances]
UserWarning: No current year days have been found for the day
2020-01-10
UserWarning: No current year days have been found for the day
2020-01-13
UserWarning: No current year days have been found for the day
2020-01-18
UserWarning: No current year days have been found for the day
2020-01-19
...

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
    best_index, test_score,
    axes) = ms.datasets_hyperparameter_validation(dataset_list, model,

```

```
hyperparameter, hyperparameter_values, time_series=True)
```

```
>>> print("Index best dataset:", str(best_index), "; t1-t2-t3:",
          str(xvalues[best_index]))
Index best dataset: 18 ; t1-t2-t3: 5-1-10

>>> print("Best polynomial grade:",
          str(datasets_best_hyperparameter_value[best_index]))
Best polynomial grade: 1

>>> print("Validation score:",
          str(datasets_train_val_score[best_index,1]))
Validation score: 0.005454695745022879

>>> print("Test score:", str(test_score))
Test score: 0.007884025168365652
```

The validation score has decreased!

Best dataset: VENTO MAX with tolerance=3

The best alternative is the one with VENTOMAX and with tolerance=3. There is an improvement as compared to the previous dataset: the validation score has decreased from 0.00580 to 0.00541

```
# `days_to_select` function
>>> days_to_select = lambda day, df, day_current_year,
                        df_current_year :
                        abs(df_current_year["VENTOMAX km/h"].loc[day] -
                        df_current_year["VENTOMAX km/h"].loc[day_current_year])<3

>>> (df_mean_best, _,
_) = tsp.add_current_year_statistics(df=df_mean_best,
df_current_year=df_mean_meteo, days_to_select=days_to_select,
columns_to_select=["mean"], y_col="mean", scale_y=True)
UserWarning: No current year days have been found for the day
2020-02-04
UserWarning: No current year days have been found for the day
2020-02-11

>>> df_mean_best
              mean    mean_1    mean_2    mean_3    mean_4
```

Datetime					
2020-01-08	38.646087	40.759729	26.099399	28.192059	54.887035
2020-01-09	39.866008	38.646087	40.759729	26.099399	28.192059
2020-01-10	63.946840	39.866008	38.646087	40.759729	26.099399
2020-01-11	53.902359	63.946840	39.866008	38.646087	40.759729
2020-01-12	42.199734	53.902359	63.946840	39.866008	38.646087
...
2020-12-27	14.997987	11.480358	14.695619	28.473515	47.568209
2020-12-28	16.317778	14.997987	11.480358	14.695619	28.473515
2020-12-29	23.536875	16.317778	14.997987	11.480358	14.695619
2020-12-30	22.759021	23.536875	16.317778	14.997987	11.480358
2020-12-31	22.005000	22.759021	23.536875	16.317778	14.997987

	mean_5	mean_6	mean_7	upTo_6_years_ago_mean
Datetime				
2020-01-08	55.216906	56.675791	76.974569	36.287863
2020-01-09	54.887035	55.216906	56.675791	35.814008
2020-01-10	28.192059	54.887035	55.216906	36.143744
2020-01-11	26.099399	28.192059	54.887035	35.806540
2020-01-12	40.759729	26.099399	28.192059	35.522678
...
2020-12-27	49.143828	43.004333	42.601218	38.902204
2020-12-28	47.568209	49.143828	43.004333	38.423588
2020-12-29	28.473515	47.568209	49.143828	37.817480
2020-12-30	14.695619	28.473515	47.568209	36.864552
2020-12-31	11.480358	14.695619	28.473515	35.864452

	6_years_ago_mean	current_year_mean
Datetime		
2020-01-08	71.802620	56.902806
2020-01-09	79.803419	53.860020
2020-01-10	71.632078	51.860875
2020-01-11	79.216124	53.371621
2020-01-12	89.111978	53.430592
...
2020-12-27	17.186082	21.778933
2020-12-28	19.039899	20.635976
2020-12-29	20.643419	20.311077
2020-12-30	19.247014	24.071164
2020-12-31	23.937203	26.838859

```
[359 rows x 11 columns]
```

Second addition

VENTOMEDIA

```
# Higher-order function, which creates the `days_to_select`
function.
# It takes in input a specific tolerance `t` and it returns the
corresponding `days_to_select` function.
>>> def create_daysToSelect_function(t):
    # `days_to_select` is a lambda function which decides whether
    # to select `day_current_year` for `day` or not.
    # It returns True if `day_current_year` has a similar
    # VENTOMEDIA with respect to `day` (i.e. the difference is
    # less than `t`)
    return (
        lambda day, df, day_current_year, df_current_year :
            abs(df_current_year["VENTOMEDIA km/h"].loc[day] -
                df_current_year["VENTOMEDIA km/h"].loc[day_current_year])<t
    )

>>> tolerances = [1, 3, 5, 10, 13, 15, 17, 20, 23] # Possible
tolerances
>>> dataset_list = [
    tsp.add_current_year_statistics(df=df_mean_best,
                                   df_current_year=df_mean_meteo,
                                   # Create `days_to_select`
                                   days_to_select = create_daysToSelect_function(t),
                                   # Only the "mean" column is of interest
                                   columns_to_select=["mean"], y_col="mean",
                                   scale_y=True)[1:]
    for t in tolerances
]
UserWarning: No current year days have been found for the day
2020-02-04
UserWarning: No current year days have been found for the day
2020-02-05
UserWarning: No current year days have been found for the day
2020-02-11
```

```
UserWarning: No current year days have been found for the day
2020-02-27
...
>>> xvalues = [str(t) for t in tolerances]

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
    best_index, test_score,
    axes) = ms.datasets_hyperparameter_validation(dataset_list,
    model, hyperparameter, hyperparameter_values, time_series=True)
```

```
>>> print("Index best dataset:", str(best_index), "; Tolerance:",
    str(xvalues[best_index]))
Index best dataset: 4 ; Tolerance: 13

>>> print("Best polynomial grade:",
    str(datasets_best_hyperparameter_value[best_index]))
Best polynomial grade: 1

>>> print("Validation score:",
    str(datasets_train_val_score[best_index,1]))
Validation score: 0.005267587773141723

>>> print("Test score:", str(test_score))
Test score: 0.006624223754847354
```

The validation score has decreased!

PRESSIONE

```
# Higher-order function, which creates the `days_to_select`
function.
# It takes in input a specific tolerance `t` and it returns the
corresponding `days_to_select` function.
>>> def create_daysToSelect_function(t):
    # `days_to_select` is a lambda function which decides whether
    # to select `day_current_year` for `day` or not.
    # It returns True if `day_current_year` has a similar
```

```

# PRESSIONE with respect to `day` (i.e. the difference is
# less than `t`)
return (
    lambda day, df, day_current_year, df_current_year :
        abs(df_current_year["PRESSIONESLM mb"].loc[day] -
            df_current_year["PRESSIONESLM mb"].loc[day_current_year])<t
)

>>> tolerances = [0.2, 0.5, 1, 3, 5, 10, 15, 20] # Possible
tolerances
>>> dataset_list = [
    tsp.add_current_year_statistics(df=df_mean_best,
                                   df_current_year=df_mean_meteo,
                                   # Create `days_to_select`
                                   days_to_select = create_daysToSelect_function(t),
                                   # Only the "mean" column is of interest
                                   columns_to_select=["mean"], y_col="mean",
                                   scale_y=True)[1:]
    for t in tolerances
]
UserWarning: No current year days have been found for the day
2020-01-09
UserWarning: No current year days have been found for the day
2020-01-10
UserWarning: No current year days have been found for the day
2020-01-11
UserWarning: No current year days have been found for the day
2020-01-13
...
>>> xvalues = [str(t) for t in tolerances]

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
    best_index, test_score,
    axes) = ms.datasets_hyperparameter_validation(dataset_list,
    model, hyperparameter, hyperparameter_values, time_series=True)

>>> print("Index best dataset:", str(best_index), "; Tolerance:",

```

```

        str(xvalues[best_index]))
Index best dataset: 3 ; Tolerance: 3

>>> print("Best polynomial grade:",
        str(datasets_best_hyperparameter_value[best_index]))
Best polynomial grade: 1

>>> print("Validation score:",
        str(datasets_train_val_score[best_index,1]))
Validation score: 0.005337474933219499

>>> print("Test score:", str(test_score))
Test score: 0.0069914173394748525

```

The validation score has decreased!

Best dataset : VENTO MED with tolerance=13

There is an improvement as compared to the previous dataset: the validation score has decreased from 0.00541 to 0.00527.

```

# `days_to_select` function
>>> days_to_select = lambda day, df, day_current_year,
    df_current_year :
    abs(df_current_year["VENTOMEDIA km/h"].loc[day] -
    df_current_year["VENTOMEDIA km/h"].loc[day_current_year])<13

>>> (df_mean_best, _,
_) = tsp.add_current_year_statistics(df=df_mean_best,
df_current_year=df_mean_meteo, days_to_select=days_to_select,
columns_to_select=["mean"], y_col="mean", scale_y=True)

>>> df_mean_best

```

	mean	mean_1	mean_2	mean_3	mean_4
Datetime					
2020-01-08	38.646087	40.759729	26.099399	28.192059	54.887035
2020-01-09	39.866008	38.646087	40.759729	26.099399	28.192059
2020-01-10	63.946840	39.866008	38.646087	40.759729	26.099399
2020-01-11	53.902359	63.946840	39.866008	38.646087	40.759729
2020-01-12	42.199734	53.902359	63.946840	39.866008	38.646087
...
2020-12-27	14.997987	11.480358	14.695619	28.473515	47.568209
2020-12-28	16.317778	14.997987	11.480358	14.695619	28.473515

```

2020-12-29 23.536875 16.317778 14.997987 11.480358 14.695619
2020-12-30 22.759021 23.536875 16.317778 14.997987 11.480358
2020-12-31 22.005000 22.759021 23.536875 16.317778 14.997987

      mean_5      mean_6      mean_7 upTo_6_years_ago_mean
Datetime
2020-01-08 55.216906 56.675791 76.974569 36.287863
2020-01-09 54.887035 55.216906 56.675791 35.814008
2020-01-10 28.192059 54.887035 55.216906 36.143744
2020-01-11 26.099399 28.192059 54.887035 35.806540
2020-01-12 40.759729 26.099399 28.192059 35.522678
...      ...      ...      ...      ...
2020-12-27 49.143828 43.004333 42.601218 38.902204
2020-12-28 47.568209 49.143828 43.004333 38.423588
2020-12-29 28.473515 47.568209 49.143828 37.817480
2020-12-30 14.695619 28.473515 47.568209 36.864552
2020-12-31 11.480358 14.695619 28.473515 35.864452

      6_years_ago_mean  current_year_mean  current_year_mean.1
Datetime
2020-01-08 71.802620 56.902806 48.400784
2020-01-09 79.803419 53.860020 47.181447
2020-01-10 71.632078 51.860875 46.368620
2020-01-11 79.216124 53.371621 48.126442
2020-01-12 89.111978 53.430592 48.651526
...      ...      ...      ...
2020-12-27 17.186082 21.778933 25.123485
2020-12-28 19.039899 20.635976 25.146538
2020-12-29 20.643419 20.311077 25.122216
2020-12-30 19.247014 24.071164 25.066962
2020-12-31 23.937203 26.838859 25.060622

[359 rows x 12 columns]

```

Third adding

PRESSIONE

```
# Higher-order function, which creates the `days_to_select`
```



```

function.
# It takes in input a specific tolerance `t` and it returns the
corresponding `days_to_select` function.
>>> def create_daysToSelect_function(t):
    # `days_to_select` is a lambda function which decides whether
    # to select `day_current_year` for `day` or not.
    # It returns True if `day_current_year` has a similar
    # PRESSIONE with respect to `day` (i.e. the difference is
    # less than `t`)
    return (
        lambda day, df, day_current_year, df_current_year :
            abs(df_current_year["PRESSIONESLM mb"].loc[day] -
                df_current_year["PRESSIONESLM mb"].loc[day_current_year])<t
    )

>>> tolerances = [0.2, 0.5, 1, 3, 5, 10, 15, 20] # Possible
tolerances
>>> dataset_list = [
    tsp.add_current_year_statistics(df=df_mean_best,
                                   df_current_year=df_mean_meteo,
                                   # Create `days_to_select`
                                   days_to_select = create_daysToSelect_function(t),
                                   # Only the "mean" column is of interest
                                   columns_to_select=["mean"], y_col="mean",
                                   scale_y=True)[1:]
    for t in tolerances
]
UserWarning: No current year days have been found for the day
2020-01-09
UserWarning: No current year days have been found for the day
2020-01-10
UserWarning: No current year days have been found for the day
2020-01-11
UserWarning: No current year days have been found for the day
2020-01-13
...
>>> xvalues = [str(t) for t in tolerances]

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

```

```
>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
      best_index, test_score,
      axes) = ms.datasets_hyperparameter_validation(dataset_list,
      model, hyperparameter, hyperparameter_values, time_series=True)
```

```
>>> print("Index best dataset:", str(best_index), "; Tolerance:",
          str(xvalues[best_index]))
```

```
Index best dataset: 0 ; Tolerance: 0.2
```

```
>>> print("Best polynomial grade:",
          str(datasets_best_hyperparameter_value[best_index]))
```

```
Best polynomial grade: 1
```

```
>>> print("Validation score:",
          str(datasets_train_val_score[best_index,1]))
```

```
Validation score: 0.005221659150174389
```

```
>>> print("Test score:", str(test_score))
```

```
Test score: 0.006919093386128838
```

The validation score has decreased!

Best dataset: PRESSIONE with tolerance=0.2

There is an improvement as compared to the previous dataset: the validation score has decreased from 0.00527 to 0.00522.

```
# `days_to_select` function
```

```
>>> days_to_select = lambda day, df, day_current_year,
      df_current_year :
      abs(df_current_year["PRESSIONESLM mb"].loc[day] -
      df_current_year["PRESSIONESLM mb"].loc[day_current_year])<0.2
```

```
>>> (df_mean_best, _,
      _) = tsp.add_current_year_statistics(df=df_mean_best,
      df_current_year=df_mean_meteo, days_to_select=days_to_select,
      columns_to_select=["mean"], y_col="mean", scale_y=True)
```

```
No current year days have been found for the day 2020-01-09
```

```
No current year days have been found for the day 2020-01-10
```

```
No current year days have been found for the day 2020-01-11
```

```
No current year days have been found for the day 2020-01-13
```

...

```
>>> df_mean_best
```

	mean	mean_1	mean_2	mean_3	mean_4
Datetime					
2020-01-08	38.646087	40.759729	26.099399	28.192059	54.887035
2020-01-09	39.866008	38.646087	40.759729	26.099399	28.192059
2020-01-10	63.946840	39.866008	38.646087	40.759729	26.099399
2020-01-11	53.902359	63.946840	39.866008	38.646087	40.759729
2020-01-12	42.199734	53.902359	63.946840	39.866008	38.646087
...
2020-12-27	14.997987	11.480358	14.695619	28.473515	47.568209
2020-12-28	16.317778	14.997987	11.480358	14.695619	28.473515
2020-12-29	23.536875	16.317778	14.997987	11.480358	14.695619
2020-12-30	22.759021	23.536875	16.317778	14.997987	11.480358
2020-12-31	22.005000	22.759021	23.536875	16.317778	14.997987

	mean_5	mean_6	mean_7	upTo_6_years_ago_mean
Datetime				
2020-01-08	55.216906	56.675791	76.974569	36.287863
2020-01-09	54.887035	55.216906	56.675791	35.814008
2020-01-10	28.192059	54.887035	55.216906	36.143744
2020-01-11	26.099399	28.192059	54.887035	35.806540
2020-01-12	40.759729	26.099399	28.192059	35.522678
...
2020-12-27	49.143828	43.004333	42.601218	38.902204
2020-12-28	47.568209	49.143828	43.004333	38.423588
2020-12-29	28.473515	47.568209	49.143828	37.817480
2020-12-30	14.695619	28.473515	47.568209	36.864552
2020-12-31	11.480358	14.695619	28.473515	35.864452

	6_years_ago_mean	current_year_mean	current_year_mean.1
Datetime			
2020-01-08	71.802620	56.902806	48.400784
2020-01-09	79.803419	53.860020	47.181447
2020-01-10	71.632078	51.860875	46.368620
2020-01-11	79.216124	53.371621	48.126442
2020-01-12	89.111978	53.430592	48.651526
...
2020-12-27	17.186082	21.778933	25.123485
2020-12-28	19.039899	20.635976	25.146538

2020-12-29	20.643419	20.311077	25.122216
2020-12-30	19.247014	24.071164	25.066962
2020-12-31	23.937203	26.838859	25.060622

current_year_mean.2

Datetime

2020-01-08	76.974569
2020-01-09	47.181447
2020-01-10	46.368620
2020-01-11	48.126442
2020-01-12	53.902359
...	...
2020-12-27	14.463473
2020-12-28	25.146538
2020-12-29	25.122216
2020-12-30	25.117861
2020-12-31	16.131590

[359 rows x 13 columns]

Fourth adding

VMEDIA VMAX PRESSIONE

```
# Higher-order function, which creates the `days_to_select` function.
# It takes in input three specific tolerances `t1-t2-t3` and it returns
the corresponding `days_to_select` function.
>>> def create_daysToSelect_function(t1, t2, t3):
    return (
        lambda day, df, day_current_year, df_current_year :
        (abs(df_current_year["VENTOMEDIA km/h"].loc[day] -
            df_current_year["VENTOMEDIA km/h"].loc[day_current_year])<t1 and
            abs(df_current_year["VENTOMAX km/h"].loc[day] -
            df_current_year["VENTOMAX km/h"].loc[day_current_year])<t2 and
            abs(df_current_year["PRESSIONESLM mb"].loc[day] -
            df_current_year["PRESSIONESLM mb"].loc[day_current_year])<t3)
    )

>>> tolerances = [1, 5, 10, 15] # Possible tolerances
```

```

>>> dataset_list = [
    tsp.add_current_year_statistics(df=df_mean_best,
                                   df_current_year=df_mean_meteo,
                                   # Create `days_to_select`
                                   days_to_select=create_daysToSelect_function(t1,t2,t3),
                                   # Only the "mean" column is of interest
                                   columns_to_select=["mean"], y_col="mean",
                                   scale_y=True)[1:]
    for t1 in tolerances for t2 in tolerances for t3 in tolerances]
UserWarning: No current year days have been found for the day
2020-01-10
UserWarning: No current year days have been found for the day
2020-01-13
UserWarning: No current year days have been found for the day
2020-01-18
UserWarning: No current year days have been found for the day
2020-01-19
...

>>> model = ms.PolynomialRegression()
>>> hyperparameter="degree"
>>> hyperparameter_values=[1,2,3]

>>> (datasets_train_val_score, datasets_best_hyperparameter_value,
    best_index, test_score,
    axes) = ms.datasets_hyperparameter_validation(dataset_list, model,
    hyperparameter, hyperparameter_values, time_series=True)

```

```

>>> print("Index best dataset:", str(best_index), "; t1-t2-t3:",
        str(xvalues[best_index]))
Index best dataset: 23 ; t1-t2-t3: 5-5-15

>>> print("Best polynomial grade:",
        str(datasets_best_hyperparameter_value[best_index]))
Best polynomial grade: 1

>>> print("Validation score:",
        str(datasets_train_val_score[best_index,1]))
Validation score: 0.00518555388716895

```

```
>>> print("Test score:", str(test_score))
Test score: 0.00724853609009831
```

The validation score has decreased!

Best dataset: tolerances 5-5-1

There is an improvement as compared to the previous dataset: the validation score has decreased from 0.00522 to 0.00518.

```
>>> days_to_select = lambda day, df, day_current_year, df_current_year :
    (abs(df_current_year["VENTOMEDIA km/h"].loc[day] -
        df_current_year["VENTOMEDIA km/h"].loc[day_current_year])<5 and
    abs(df_current_year["VENTOMAX km/h"].loc[day] -
        df_current_year["VENTOMAX km/h"].loc[day_current_year])<5 and
    abs(df_current_year["PRESSIONESLM mb"].loc[day] -
        df_current_year["PRESSIONESLM mb"].loc[day_current_year])<15)
```

```
>>> (df_mean_best, X,
    y) = tsp.add_current_year_statistics(df=df_mean_best,
    df_current_year=df_mean_meteo, days_to_select=days_to_select,
    columns_to_select=["mean"], y_col="mean", scale_y=True)
```

```
>>> df_mean_best
```

	mean	mean_1	mean_2	mean_3	mean_4
Datetime					
2020-01-08	38.646087	40.759729	26.099399	28.192059	54.887035
2020-01-09	39.866008	38.646087	40.759729	26.099399	28.192059
2020-01-10	63.946840	39.866008	38.646087	40.759729	26.099399
2020-01-11	53.902359	63.946840	39.866008	38.646087	40.759729
2020-01-12	42.199734	53.902359	63.946840	39.866008	38.646087
...
2020-12-27	14.997987	11.480358	14.695619	28.473515	47.568209
2020-12-28	16.317778	14.997987	11.480358	14.695619	28.473515
2020-12-29	23.536875	16.317778	14.997987	11.480358	14.695619
2020-12-30	22.759021	23.536875	16.317778	14.997987	11.480358
2020-12-31	22.005000	22.759021	23.536875	16.317778	14.997987

	mean_5	mean_6	mean_7	upTo_6_years_ago_mean
Datetime				
2020-01-08	55.216906	56.675791	76.974569	36.287863
2020-01-09	54.887035	55.216906	56.675791	35.814008
2020-01-10	28.192059	54.887035	55.216906	36.143744
2020-01-11	26.099399	28.192059	54.887035	35.806540

2020-01-12	40.759729	26.099399	28.192059	35.522678
...
2020-12-27	49.143828	43.004333	42.601218	38.902204
2020-12-28	47.568209	49.143828	43.004333	38.423588
2020-12-29	28.473515	47.568209	49.143828	37.817480
2020-12-30	14.695619	28.473515	47.568209	36.864552
2020-12-31	11.480358	14.695619	28.473515	35.864452

	6_years_ago_mean	current_year_mean	current_year_mean.1
Datetime			
2020-01-08	71.802620	56.902806	48.400784
2020-01-09	79.803419	53.860020	47.181447
2020-01-10	71.632078	51.860875	46.368620
2020-01-11	79.216124	53.371621	48.126442
2020-01-12	89.111978	53.430592	48.651526
...
2020-12-27	17.186082	21.778933	25.123485
2020-12-28	19.039899	20.635976	25.146538
2020-12-29	20.643419	20.311077	25.122216
2020-12-30	19.247014	24.071164	25.066962
2020-12-31	23.937203	26.838859	25.060622

	current_year_mean.2	current_year_mean.3
Datetime		
2020-01-08	76.974569	51.768905
2020-01-09	47.181447	53.860020
2020-01-10	46.368620	51.860875
2020-01-11	48.126442	50.341374
2020-01-12	53.902359	50.697472
...
2020-12-27	14.463473	19.177966
2020-12-28	25.146538	17.899923
2020-12-29	25.122216	18.574172
2020-12-30	25.117861	20.902818
2020-12-31	16.131590	24.139165

[359 rows x 14 columns]

Conclusion

Now the results of the previous case study are discussed.

On the whole, the validation score has dramatically decreased. It has decreased from 0.01523 of the trivial model to 0.00518.

The final test score is 0.00725.

The best dataset which has been selected contains 14 features. More specifically, it contains the response feature “mean” and 13 explanatory features.

```
>>> df_mean_best.info()
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 359 entries, 2020-01-08 to 2020-12-31
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype  
---  -
 0   mean                                359 non-null    float64
 1   mean_1                              359 non-null    float64
 2   mean_2                              359 non-null    float64
 3   mean_3                              359 non-null    float64
 4   mean_4                              359 non-null    float64
 5   mean_5                              359 non-null    float64
 6   mean_6                              359 non-null    float64
 7   mean_7                              359 non-null    float64
 8   upTo_6_years_ago_mean               359 non-null    float64
 9   6_years_ago_mean                    359 non-null    float64
10   current_year_mean                   359 non-null    float64
11   current_year_mean.1                 359 non-null    float64
12   current_year_mean.2                 359 non-null    float64
13   current_year_mean.3                 359 non-null    float64
dtypes: float64(14)
memory usage: 50.2 KB

>>> df_mean_best
          mean    mean_1    mean_2    mean_3    mean_4
Datetime
2020-01-08  38.646087  40.759729  26.099399  28.192059  54.887035
2020-01-09  39.866008  38.646087  40.759729  26.099399  28.192059
2020-01-10  63.946840  39.866008  38.646087  40.759729  26.099399
2020-01-11  53.902359  63.946840  39.866008  38.646087  40.759729
2020-01-12  42.199734  53.902359  63.946840  39.866008  38.646087
```


...
2020-12-27	14.997987	11.480358	14.695619	28.473515	47.568209
2020-12-28	16.317778	14.997987	11.480358	14.695619	28.473515
2020-12-29	23.536875	16.317778	14.997987	11.480358	14.695619
2020-12-30	22.759021	23.536875	16.317778	14.997987	11.480358
2020-12-31	22.005000	22.759021	23.536875	16.317778	14.997987

	mean_5	mean_6	mean_7	upTo_6_years_ago_mean
Datetime				
2020-01-08	55.216906	56.675791	76.974569	36.287863
2020-01-09	54.887035	55.216906	56.675791	35.814008
2020-01-10	28.192059	54.887035	55.216906	36.143744
2020-01-11	26.099399	28.192059	54.887035	35.806540
2020-01-12	40.759729	26.099399	28.192059	35.522678
...
2020-12-27	49.143828	43.004333	42.601218	38.902204
2020-12-28	47.568209	49.143828	43.004333	38.423588
2020-12-29	28.473515	47.568209	49.143828	37.817480
2020-12-30	14.695619	28.473515	47.568209	36.864552
2020-12-31	11.480358	14.695619	28.473515	35.864452

	6_years_ago_mean	current_year_mean	current_year_mean.1
Datetime			
2020-01-08	71.802620	56.902806	48.400784
2020-01-09	79.803419	53.860020	47.181447
2020-01-10	71.632078	51.860875	46.368620
2020-01-11	79.216124	53.371621	48.126442
2020-01-12	89.111978	53.430592	48.651526
...
2020-12-27	17.186082	21.778933	25.123485
2020-12-28	19.039899	20.635976	25.146538
2020-12-29	20.643419	20.311077	25.122216
2020-12-30	19.247014	24.071164	25.066962
2020-12-31	23.937203	26.838859	25.060622

	current_year_mean.2	current_year_mean.3
Datetime		
2020-01-08	76.974569	51.768905
2020-01-09	47.181447	53.860020
2020-01-10	46.368620	51.860875
2020-01-11	48.126442	50.341374

2020-01-12	53.902359	50.697472
...
2020-12-27	14.463473	19.177966
2020-12-28	25.146538	17.899923
2020-12-29	25.122216	18.574172
2020-12-30	25.117861	20.902818
2020-12-31	16.131590	24.139165


[359 rows x 14 columns]

The best model which has been selected is PolynomialRegression with degree=1, i.e. the linear regression model.

```
>>> model = ms.PolynomialRegression(degree=1)
```

To conclude, the plots which describe the goodness of the predictions are shown. The improvement of the goodness of predictions as compared to the first model is graphically visible.

```
>>> ms.plot_predictions(X ,y ,model, xvalues=df_mean_best.index,
xlabel="Days")
<AxesSubplot:title={'center': 'Actual vs Predicted values'},
xlabel='Days'>
```



```
>>> ms.plot_predictions(X ,y ,model, plot_type=1,
xvalues=df_mean_best.index, xlabel="Days")
<AxesSubplot:title={'center': 'Actual vs Predicted values'},
xlabel='Actual values', ylabel='Predicted values'>
```

